

Section 8

Matching and Weighting Estimators

Sooahn Shin

GOV 2003

Nov 4, 2021

Overview

- Logistics:
 - **Pset 7 released!** Due at 11:59 pm (ET) on Nov 10
- Today's topics:
 1. Matching estimators
 2. Weighting estimators

[Review] Observational studies

- **Identification**

- Most common observational assumptions:
 - No unmeasured confounders: $D_i \perp\!\!\!\perp (Y_i(0), Y_i(1)) \mid \mathbf{X}_i$
 - Overlap/positivity: $0 < \mathbb{P}(D_i = 1 \mid \mathbf{X}_i = \mathbf{x}) < 1$
- Estimand:
 - ATE = $\mathbb{E}[Y_i(1) - Y_i(0)]$ (We identified this in Module 5)
 - ATT = $\mathbb{E}[Y_i(1) - Y_i(0) \mid D_i = 1]$
 - ATC = $\mathbb{E}[Y_i(1) - Y_i(0) \mid D_i = 0]$

[Review] Observational studies

- **Identification**

- Most common observational assumptions:
 - No unmeasured confounders: $D_i \perp\!\!\!\perp (Y_i(0), Y_i(1)) \mid \mathbf{X}_i$
 - Overlap/positivity: $0 < \mathbb{P}(D_i = 1 \mid \mathbf{X}_i = \mathbf{x}) < 1$
- Estimand:
 - ATE = $\mathbb{E}[Y_i(1) - Y_i(0)]$ (We identified this in Module 5)
 - ATT = $\mathbb{E}[Y_i(1) - Y_i(0) \mid D_i = 1]$
 - ATC = $\mathbb{E}[Y_i(1) - Y_i(0) \mid D_i = 0]$

- **Estimation**

- Regression estimators: $\hat{\mu}_1(\mathbf{x})$ and $\hat{\mu}_0(\mathbf{x})$ (Module 5)
- Matching estimator (for ATT):

$$\hat{\tau}_m = \frac{1}{n_1} \sum_{i=1}^n D_i \left(Y_i - \frac{1}{M} \sum_{j \in \mathcal{J}_M(i)} Y_j \right)$$

- Weighting estimators:
 - Horvitz-Thompson estimator (= IPW estimator)
 - Hajek estimator (normalized weights)

Types of matching

- Exact matching: choose matches that have the same value of \mathbf{X}_i
 - Cf: Coarsened Exact Matching (CEM)

Types of matching

- Exact matching: choose matches that have the same value of \mathbf{X}_i
 - Cf: Coarsened Exact Matching (CEM)
- Mahalanobis distance matching: use distance metrics in case of high dimensional \mathbf{X}_i

Types of matching

- Exact matching: choose matches that have the same value of \mathbf{X}_i
 - Cf: Coarsened Exact Matching (CEM)
- Mahalanobis distance matching: use distance metrics in case of high dimensional \mathbf{X}_i
- Propensity score matching:

$$(Y_i(0), Y_i(1)) \perp\!\!\!\perp D_i \mid \mathbf{X}_i \Rightarrow (Y_i(0), Y_i(1)) \perp\!\!\!\perp D_i \mid \pi(\mathbf{X}_i)$$

- This holds under **true** propensity score $\pi(\mathbf{X}_i)$.
- We need to estimate it ($\hat{\pi}(\mathbf{X}_i)$): e.g., using logistic regression (can add interactions) or machine learning.
- Have to check if \mathbf{X}_i is actually balanced.

Types of matching

- Other choices
 - Matching ratio: m -to-one matching
 - w/ or w/o replacement: consider the number of control units
 - Caliper: drop poor matches (estimand changes)

Types of matching

- Other choices
 - Matching ratio: m -to-one matching
 - w/ or w/o replacement: consider the number of control units
 - Caliper: drop poor matches (estimand changes)
- Algorithm:
 - Greedy algorithm: pair two units with the shortest distance, set them aside, and repeat \leadsto depends on order and thus may not be optimal
 - Optimal matching:
 - \mathbf{D} : $n \times n$ matrix of pairwise distance or a cost matrix
 - Select n elements of \mathbf{D} such that there is only one element in each row and one element in each column and the sum of pairwise distances is minimized \leadsto linear sum assignment problem

Types of matching

- Other choices
 - Matching ratio: m -to-one matching
 - w/ or w/o replacement: consider the number of control units
 - Caliper: drop poor matches (estimand changes)
- Algorithm:
 - Greedy algorithm: pair two units with the shortest distance, set them aside, and repeat \leadsto depends on order and thus may not be optimal
 - Optimal matching:
 - \mathbf{D} : $n \times n$ matrix of pairwise distance or a cost matrix
 - Select n elements of \mathbf{D} such that there is only one element in each row and one element in each column and the sum of pairwise distances is minimized \leadsto linear sum assignment problem
- Assessing balance
 - standardized mean differences
 - Kolmogorov–Smirnov statistic (comparing distributions)

Matching estimators

- Workflow (in general):
 1. Check the balance before the matching
 2. Choose matching type (compute/estimate the distance/balancing score if necessary)
 3. Conduct matching (check the matched dataset)
 4. Check the balance after the matching
 5. Estimate ATT using matching estimator ($\hat{\tau}_m$)
 6. Estimate the standard error
 - w/o replacement: cluster bootstrap
 - w/ replacement: use Abadie and Imbens (2006) estimator

Matching estimators

- Workflow (in general):
 1. Check the balance before the matching
 2. Choose matching type (compute/estimate the distance/balancing score if necessary)
 3. Conduct matching (check the matched dataset)
 4. Check the balance after the matching
 5. Estimate ATT using matching estimator ($\hat{\tau}_m$)
 6. Estimate the standard error
 - w/o replacement: cluster bootstrap
 - w/ replacement: use Abadie and Imbens (2006) estimator
- Useful packages:
 - `cobalt`: for balance check (`bal.tab()` and `love.plot()`)
 - `MatchIt`: for matching
 - `Matching`: for matching + estimating
 - Machine learning packages for estimating $\hat{\pi}(\mathbf{X}_i)$: e.g., `randomForest`
 - Optimal matching: `clue::solve_LSAP()`

Example: LaLonde dataset

- The effectiveness of a job training program (National Supported Work Demonstration; NSW) on wage increases.
- The federal government instituted a randomized evaluation of this program
- How well the result may be recovered when the experimental controls are replaced with a set of observational controls (Population Survey of Income Dynamics; PSID)?
- **Problem:** Imbalances between the experimental and observational data \leadsto use matching

Example: LaLonde dataset

- Data:
 - Treated: 185 units from NSW
 - Control: 2490 units from PSID
 - Treatment: Participation in the job training program (`nsw`)
 - Outcome: 1978 earnings (in dollars; `re78`)
 - Pre-treatment covariates: age, race, marriage, past earnings, past employment

Example: Balance before matching

```
library(cobalt)
bal.tab(x = dat[,pretreat_covariates],
        treat = dat$nsu, continuous = "std", binary = "std")
```

Example: Balance before matching

```
## Balance Measures
##           Type Diff.Un
## age      Contin. -1.0094
## educ      Contin. -0.6805
## black     Binary  1.4816
## hisp      Binary  0.1288
## married   Binary -1.8453
## re74      Contin. -1.7178
## re75      Contin. -1.7744
## u74       Binary  1.6454
## u75       Binary  1.2309
##
## Sample sizes
##      Control Treated
## All      2490      185
```


Example: Propensity score matching


```
# Estimate propensity score using logistic regression  
# Propensity score  
pscores <- glm(nsw ~ age + I(age^2) + black + married + hisp + u74,  
               family = binomial(), data = dat)$fitted.values  
# Conduct one-to-one nearest neighbor propensity score matching  
library(Matching)  
match_ps <- Match(Y=dat$re78, Tr=dat$nsw,  
                 X=pscores, M=1, replace = TRUE, ties = FALSE)  
summary(match_ps)
```

Example: Propensity score matching

```
##
## Estimate... -941.15
## SE..... 903.91
## T-stat..... -1.0412
## p.val..... 0.29779
##
## Original number of observations..... 2675
## Original number of treated obs..... 185
## Matched number of observations..... 185
## Matched number of observations (unweighted). 185
```


Example: Propensity score matching

Propensity Score Distribution Before Matching



Example: Propensity score matching


Propensity Score Distribution After Matching



Example: Propensity score matching

```
library(cobalt)
love.plot(nsw ~ age+educ+black+hispanic+married+re74+re75+u74+u75,
  data = dat,
  stats = "mean.diffs",
  weights = data.frame(Matched = get.w(match_ps)),
  method = c("matching"), binary = "std")
```

Example: Propensity score matching



Weighting estimators

- Matching is actually a special case of a weighting estimator
- Horvitz-Thompson estimator: weight by inverse propensity score.



$$\widehat{ATE} = \widehat{\tau}_{ipw} = \frac{1}{n} \sum_{i=1}^n \left(\frac{D_i Y_i}{\widehat{\pi}(\mathbf{X}_i)} - \frac{(1 - D_i) Y_i}{1 - \widehat{\pi}(\mathbf{X}_i)} \right)$$

- Would be unbiased if we knew the true propensity scores, $\pi(\mathbf{X}_i)$ (Pset 7 Q2 Bonus)
- Under no unmeasured confounders, $\widehat{\tau}_{ipw} \xrightarrow{P} \tau$ (consistent)
- Hajek estimator: normalizes the weights
- Potential of extreme weights due to lack of overlap: $\pi(\mathbf{X}_i)$ close to 0 or 1
 - Winsorizing: trim weights beyond 5th and 95th percentile

Example

Generating propensity score weights for the ATT

```
W.out <- WeightIt::weightit(nsw ~ age + I(age^2) + black + married + hisp +  
  method = "ps", estimand = "ATT")
```



Estimating ATT with Weights

```
## append estimated weights
```

```
dat <- dat %>% mutate(weights = W.out$weights)
```

```
## ATT
```

```
att_ipw <- function(dat, indices = NULL) {  
  if (is.null(indices)) indices <- 1:nrow(dat)  
  dat <- dat[indices,]
```

```
  weights <- dat %>% filter(treat == 0) %>% pull(weights)
```

```
  reweights <- weights / sum(weights)
```

```
  Y1 <- dat %>% filter(treat == 1) %>% pull(re78)
```

```
  Y0 <- dat %>% filter(treat == 0) %>% pull(re78)
```

```
  att_ht <- sum(Y1 - Y0 * weights) / nobs
```

```
  att_hjk <- mean(Y1) - sum(reweights * Y0)
```

```
  return(c(att_ht, att_hjk))
```

```
}
```

```
## Use bootstrap for estimating SE
```